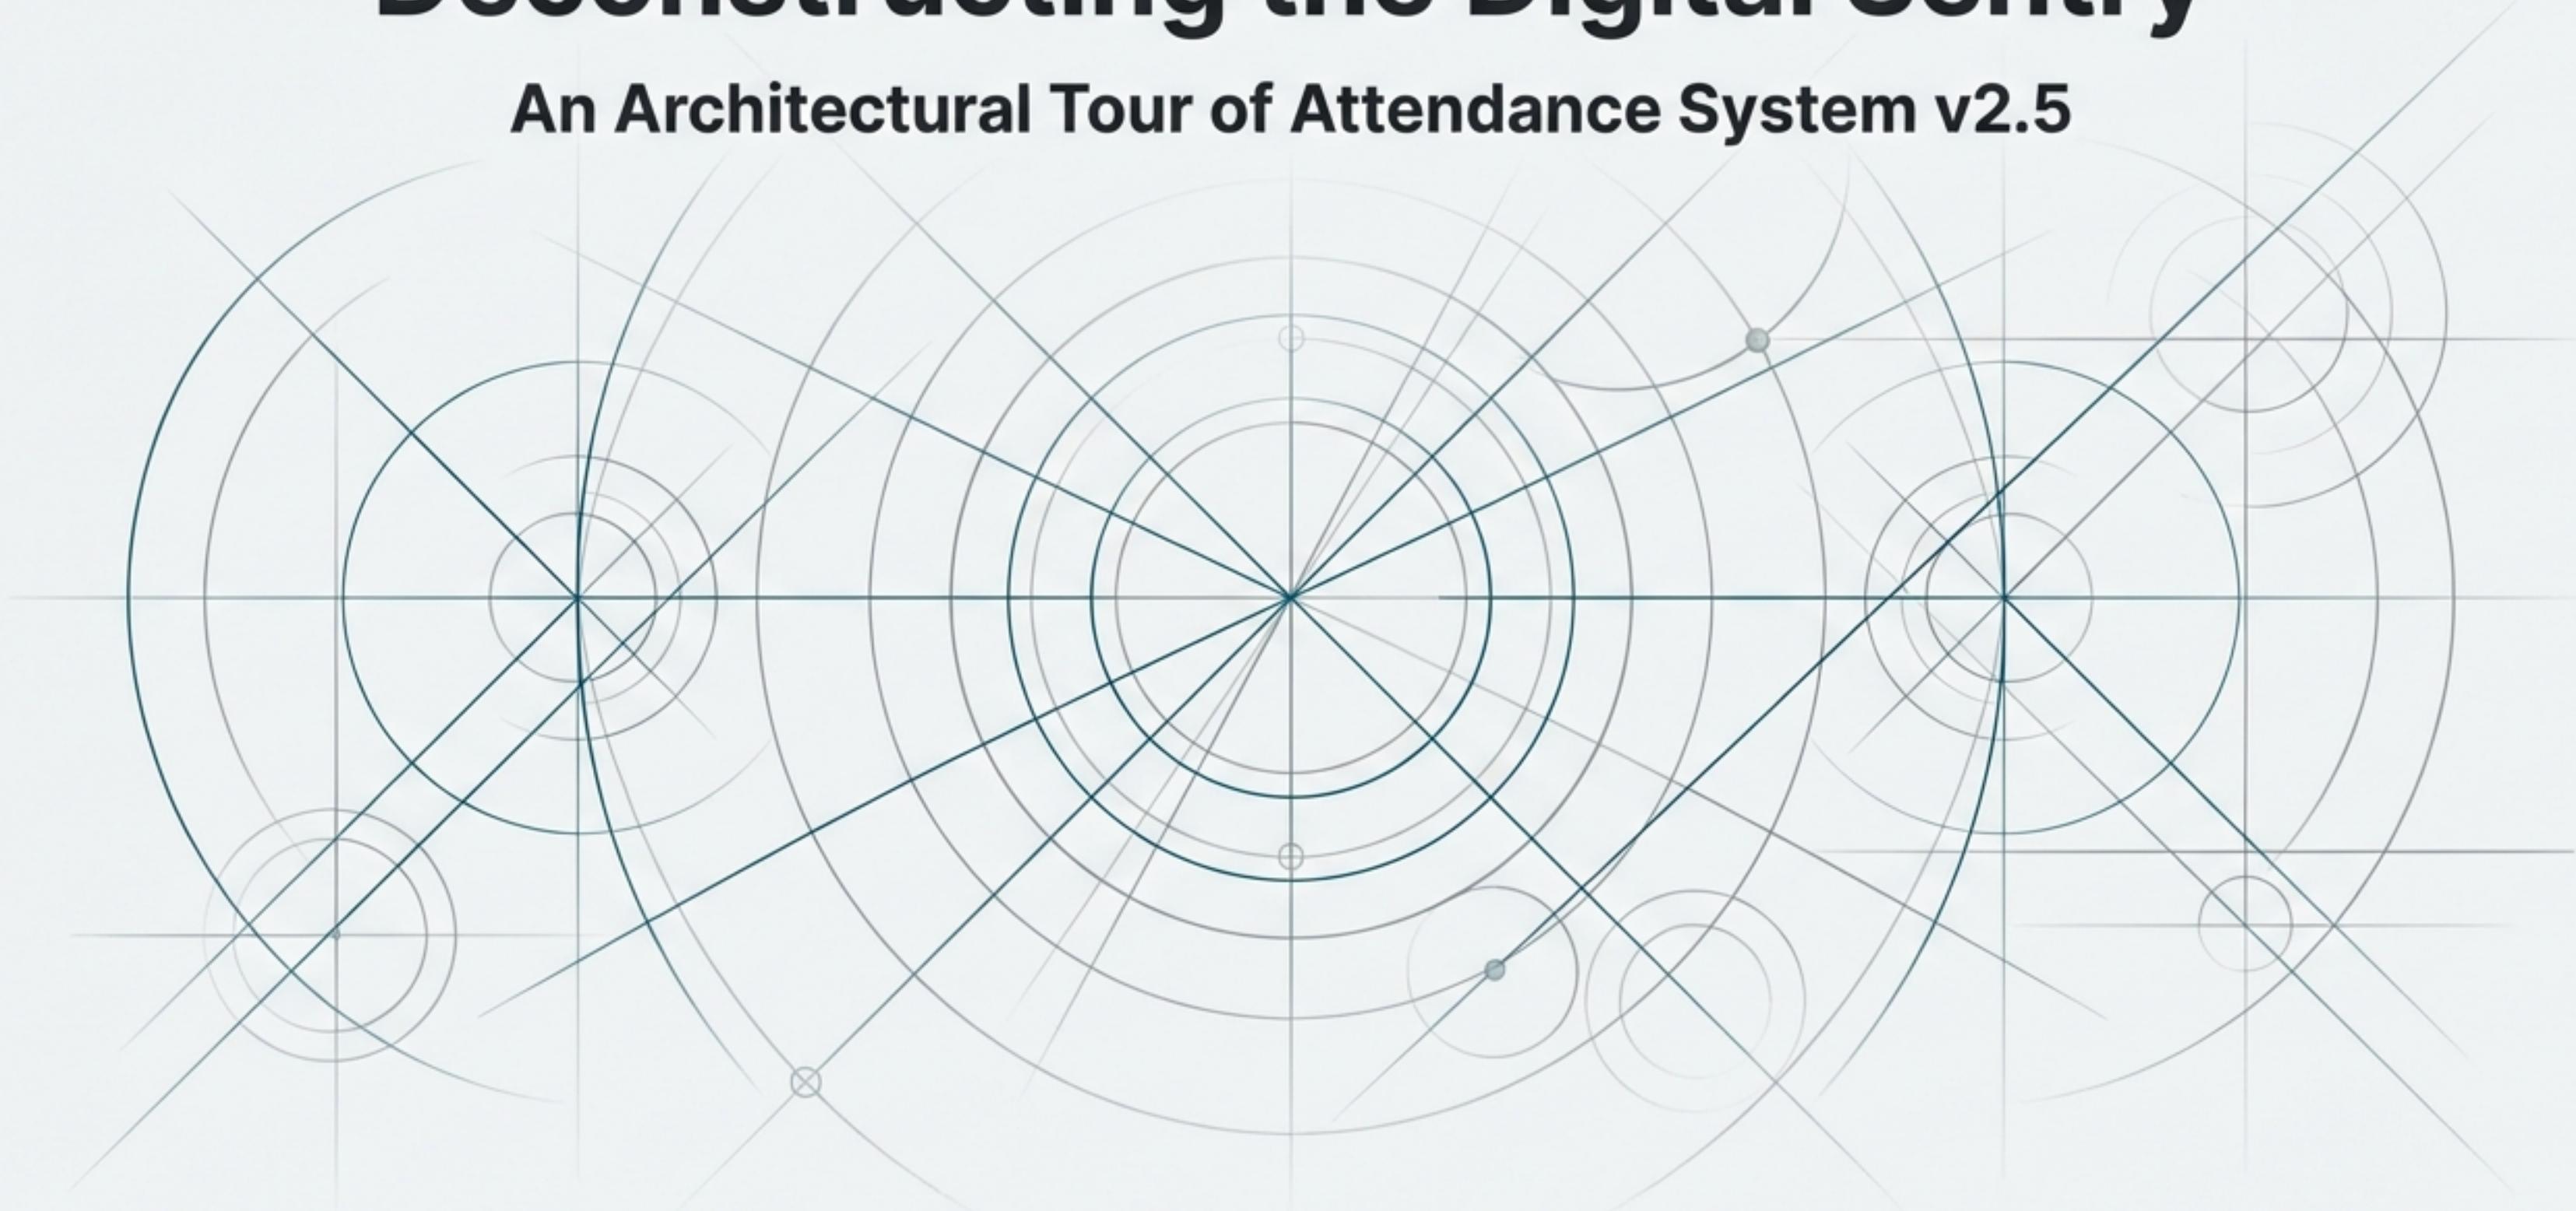


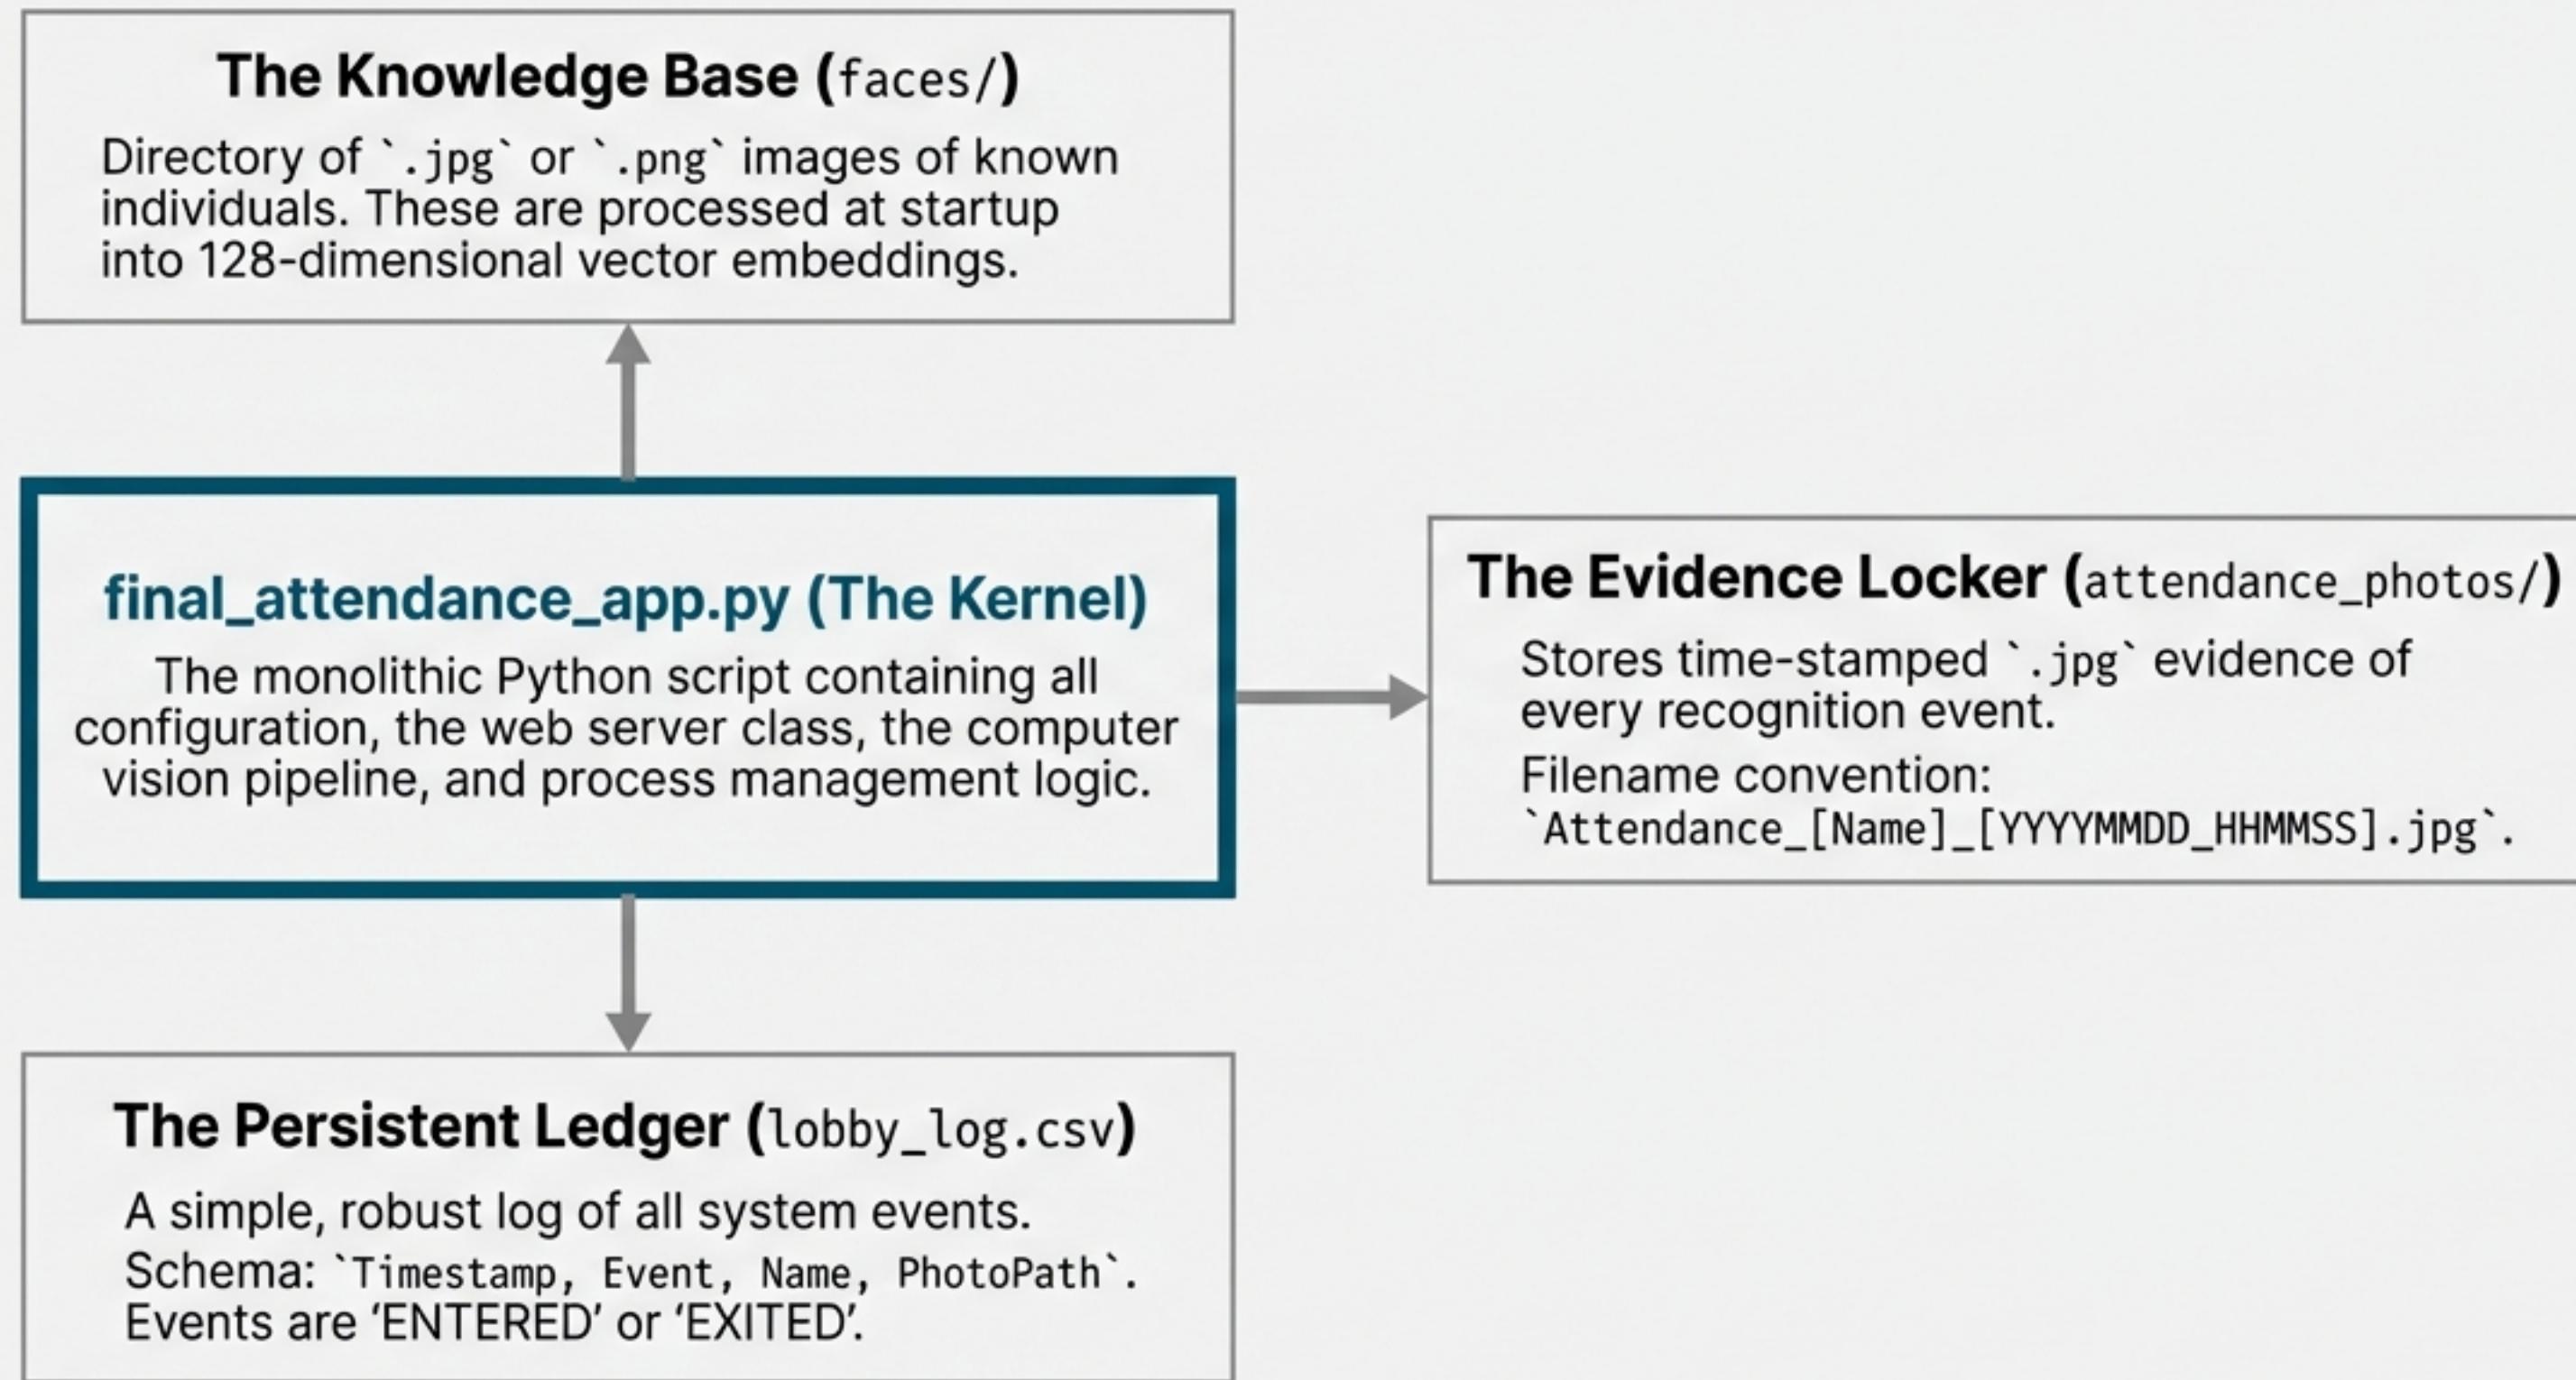
# Deconstructing the Digital Sentry

## An Architectural Tour of Attendance System v2.5



A “Teeth to Bone” analysis of the system’s design, logic, and core optimizations.

# The System Blueprint: Anatomy of a Monolithic Kernel



# The System's State: Core Variables in Memory

`known_face_encodings` (List of Lists of Floats)

## The Identity Library

The mathematical representation of every known face.  
A list of 128-d vectors.

`known_face_names` (List of Strings)

## The Name Roster

A parallel list mapping each vector to a human-readable name. `encodings[i]` corresponds to `names[i]`.

`present_people` (Dictionary)

## The Live Tracker

The system's short-term memory. **Key:** Name, **Value:** Last Seen Timestamp. Prevents log duplication and detects exits.

`process_this_frame` (Integer)

## The CPU Governor

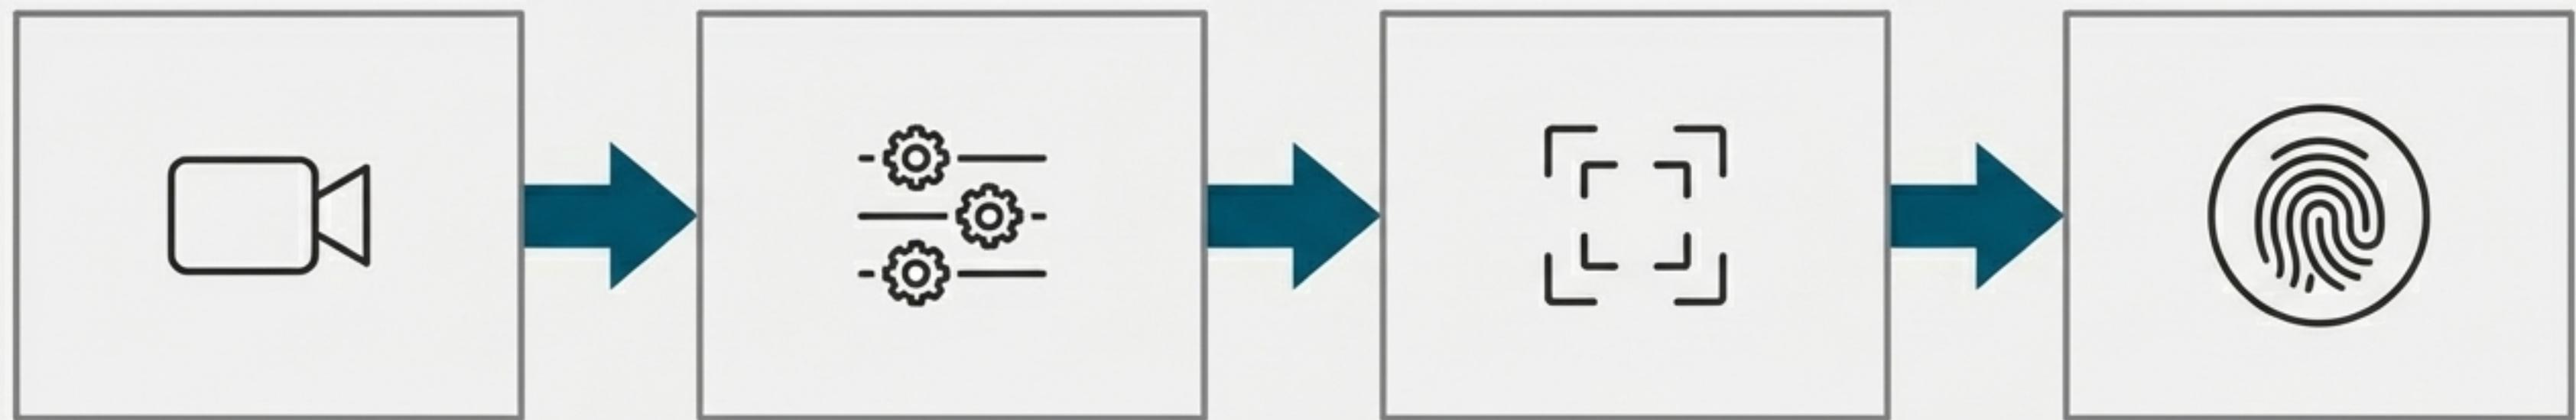
A simple counter used with modulo arithmetic to ensure the computationally expensive recognition logic only runs on select frames.

`server_process` (multiprocessing.Process)

## The Process Handle

A handle to the child process running the Flask web server. Is `None` when the server is off, allowing for state tracking.

# Inside the Brain: The Computer Vision Pipeline



## Acquisition

Capture raw video frame from the source.

## Pre-processing

Resize and convert the frame for optimal performance.

## Detection

Locate all potential faces within the frame ('Where').

## Recognition

Identify the located faces against the knowledge base ('Who').

This pipeline runs in a continuous `while True` loop, processing the video stream in real-time.

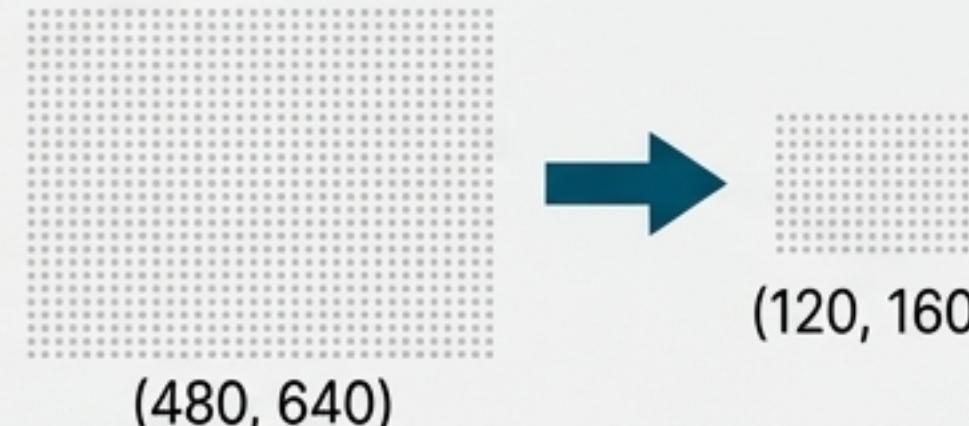
# Step 1 & 2: Frame Acquisition and Strategic Optimization

## Acquisition: The Raw Data

- Source: `cv2.VideoCapture(0)` (Default Webcam)
- Data Structure: A numpy array of shape `(480, 640, 3)`
- Total Pixels: 307,200

## Pre-processing: The Rationale for Speed

### Resizing



Code: `cv2.resize(frame, (0, 0), fx=0.25, fy=0.25)`

"Input": (480, 640) -> "Output": (120, 160)

"Pixel Reduction": 307,200 -> 19,200 (A 94% reduction)

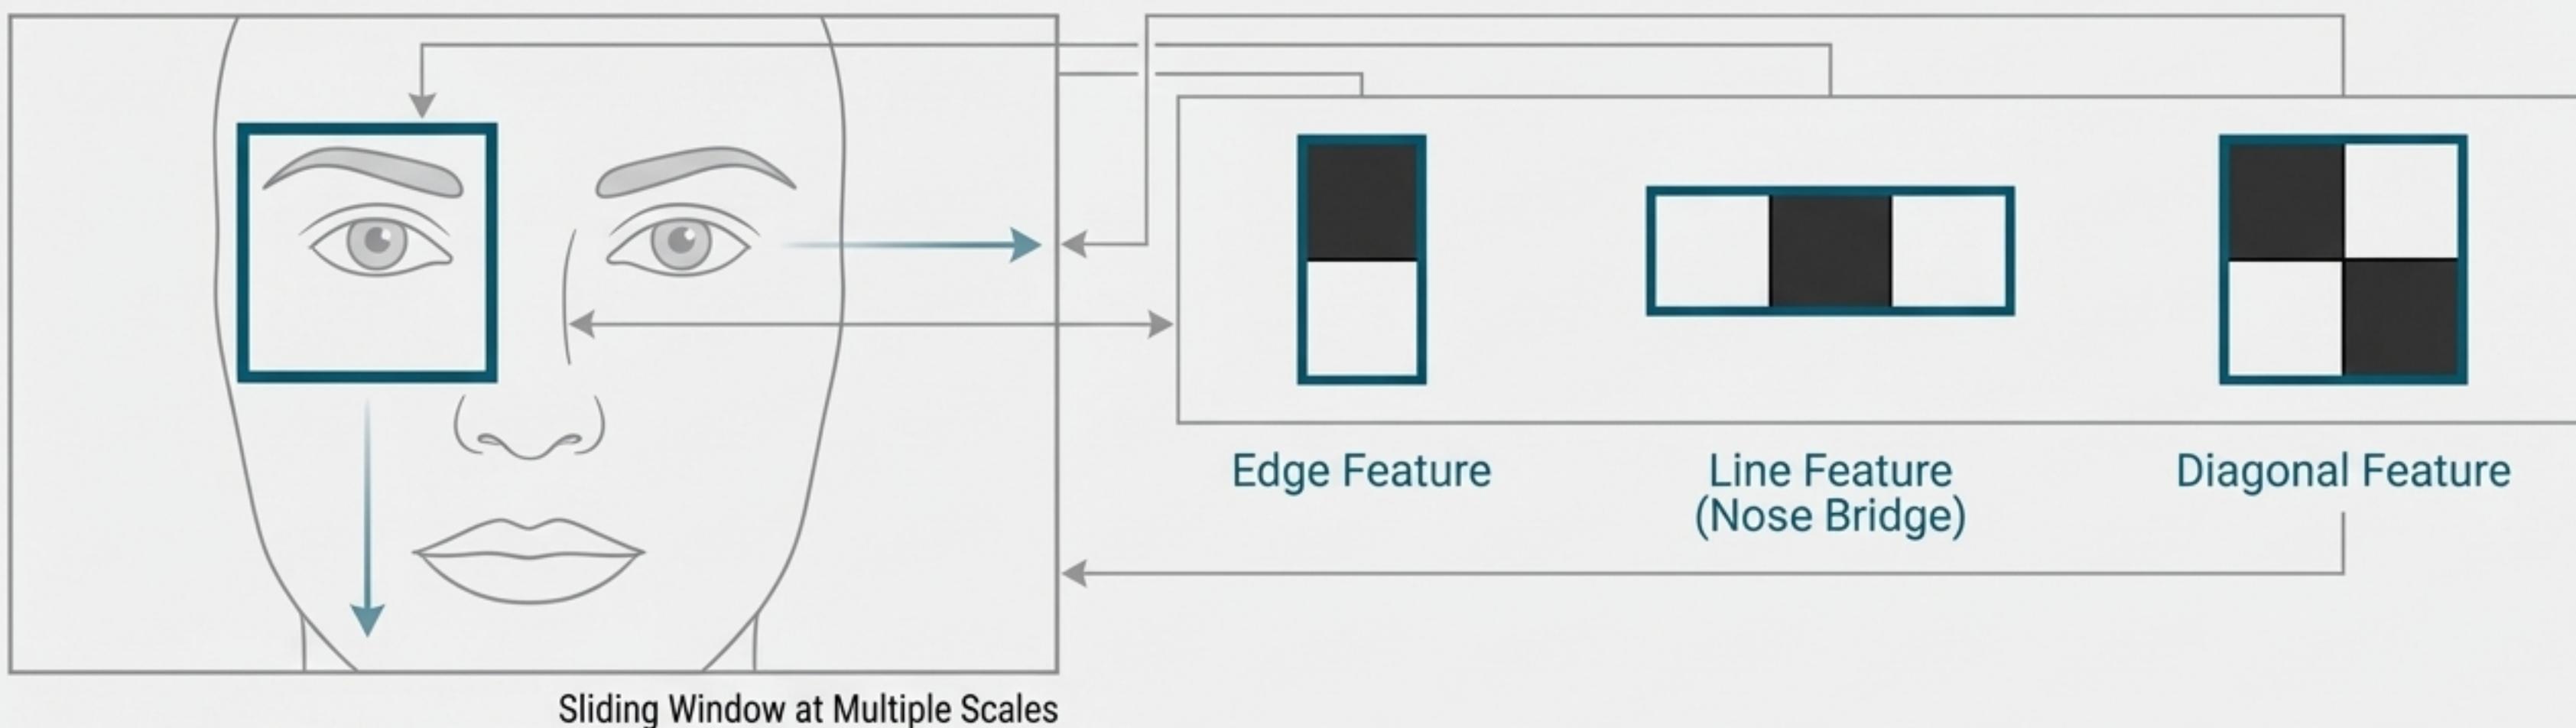
**WHY:** Reduces data dimensionality before it hits detection algorithms, which have  $O(N^2)$  or  $O(N \log N)$  complexity. This is the single most important performance tweak.

### Color Conversion

The frame is converted to both Grayscale (for Haar detection) and RGB (for dlib recognition) to meet the specific requirements of each library.

# Step 3: Detection - Answering “Where is the Face?”

The Viola-Jones Framework with Haar Cascades



## Description of Mechanism

1. The algorithm slides a window across the grayscale image at multiple scales.
2. At each position, it looks for simple, pre-defined "Haar-like features" (e.g., a light region surrounded by two dark regions, characteristic of eyes).
3. A cascade of checks rapidly discards non-face regions.

**Key Settings:** `scaleFactor=1.1`, `minNeighbors=5`

**The Output:** The result is not an identity, but simply a list of rectangles  $[(x, y, w, h)]$ , defining the coordinates of every face found in the frame.

# Step 4: Recognition - Answering ‘Who is this Person?’

## Deep Learning with dlib’s ResNet Model

### The Gate

The system conserves resources by only running this heavy process on a subset of frames.

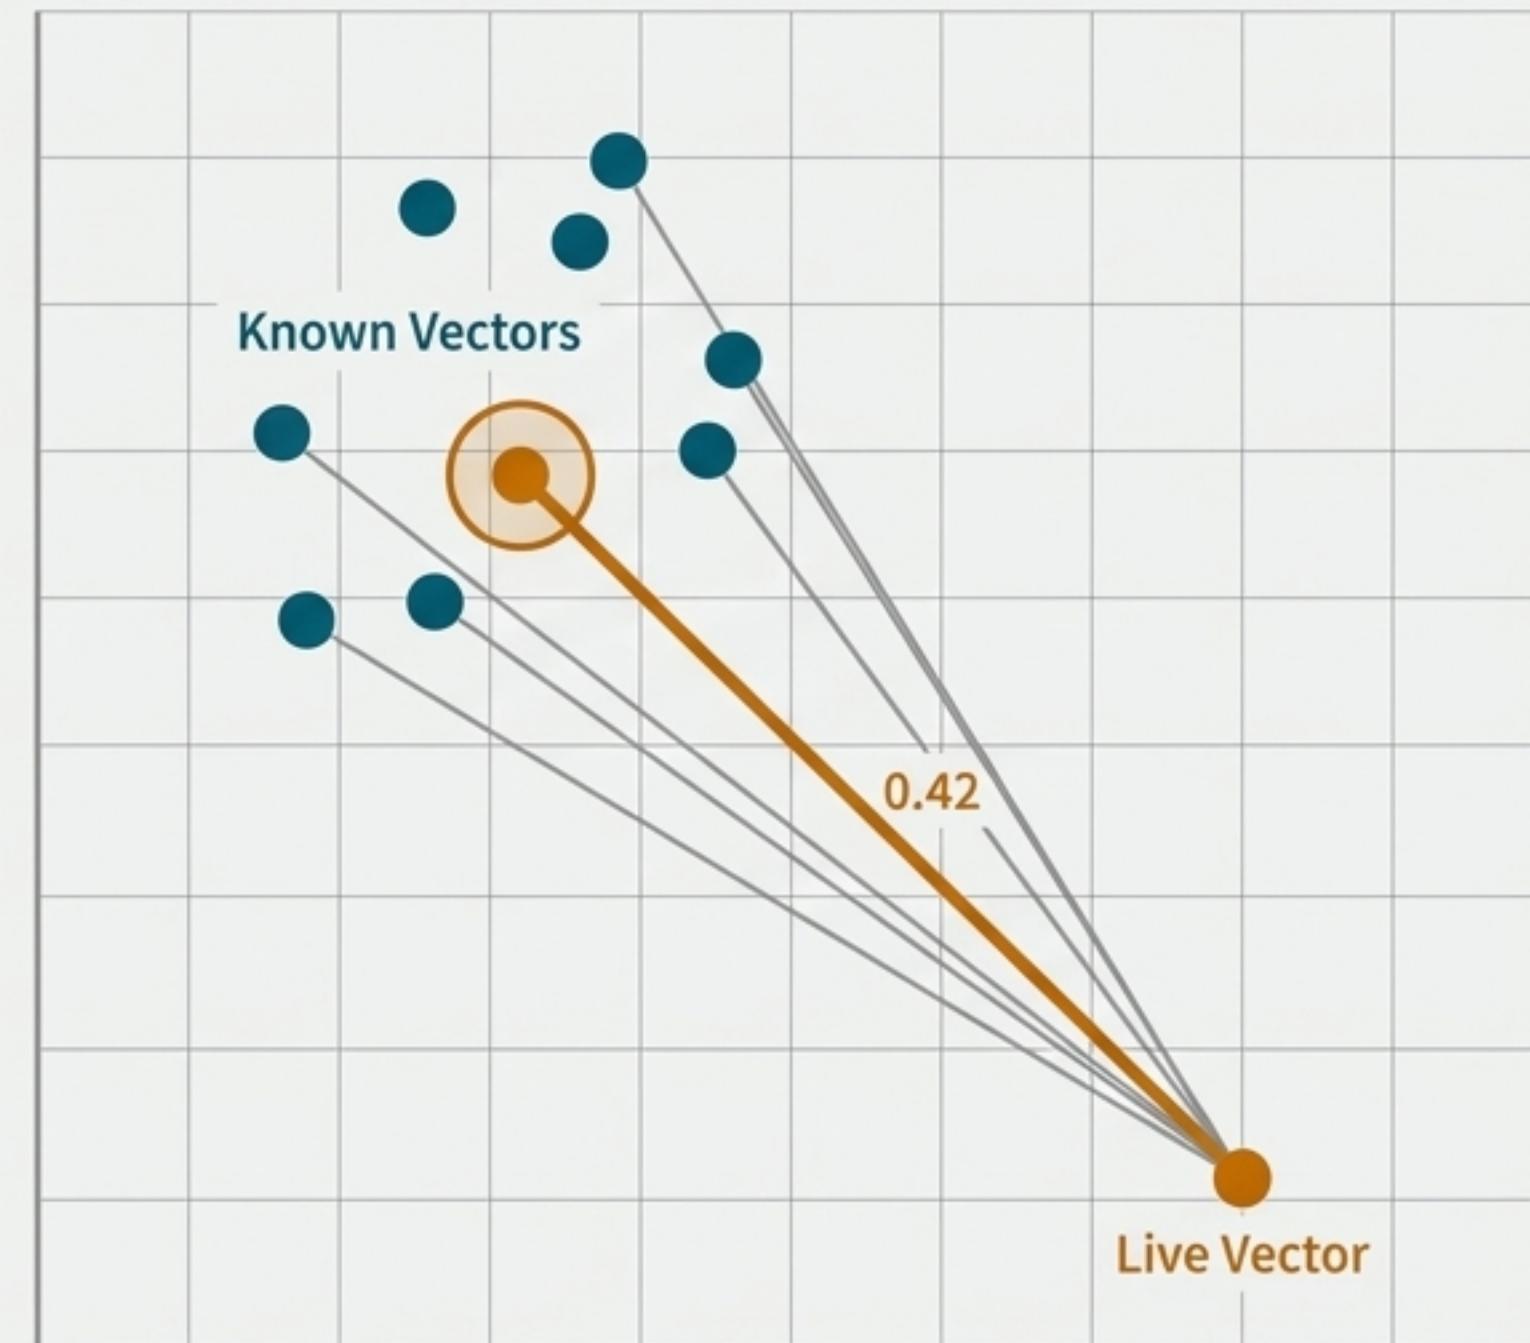
Condition: `process_this_frame % 5 == 0`

### Encoding to a Vector

The `160x120` RGB image and face location rectangle are passed to dlib. A Deep Residual Network (ResNet) model processes the face and outputs a 128-dimensional vector (an ‘embedding’). This vector is a unique mathematical signature for that face.

### Matching by Distance

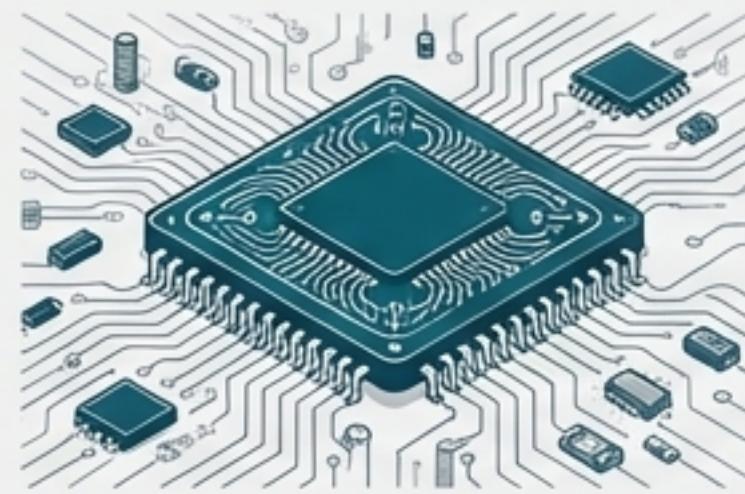
The new vector is compared against all vectors in `known\_face\_encodings` by calculating the Euclidean distance. The name associated with the vector having the smallest distance is chosen as the match, provided the distance is below the tolerance threshold of 0.6.



Distance < 0.6 indicates a positive match.

# A Masterclass in Memory Architecture: Heavy vs. Light Processes

## Main Process: The Vision Engine



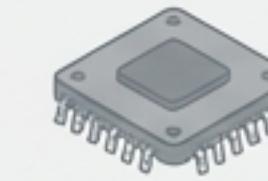
**~600MB**

### Loads:

- 💻 `cv2` (OpenCV)
- 📦 `numpy`
- 📦 `face\_recognition` (dlib wrapper)
- (chip) `dlib` (C++ Engine)

The dlib ResNet and landmark predictor models account for ~500MB of the total footprint.

## Child Process: The Web Interface



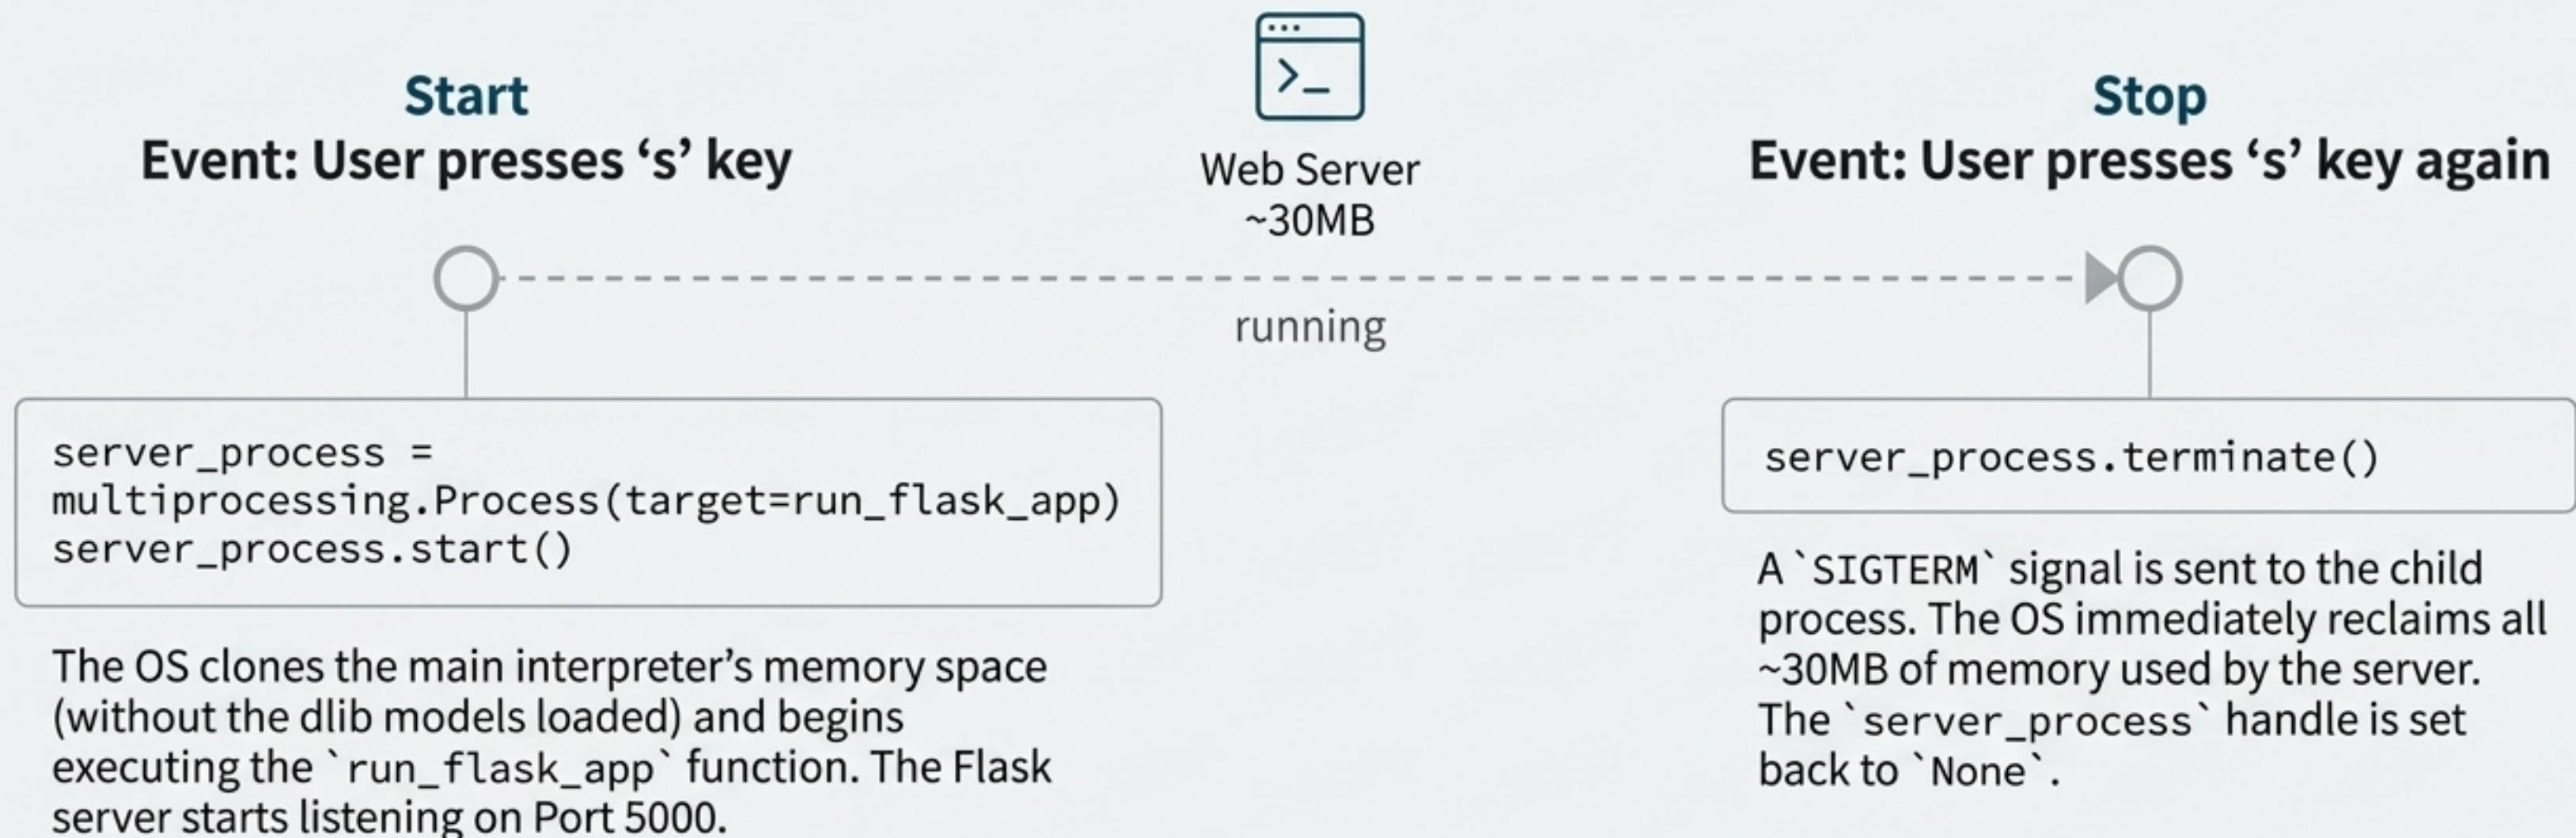
**~30MB**

### Loads:

- banana `flask`
- terminal `socket`
- stack `threading`

The `import face\_recognition` statement is scoped *\*inside\** the main vision loop. Because the web server is spawned as a separate process from a different function, it *\*never\** imports dlib or its models, resulting in a 20x smaller memory footprint.

# The On-Demand Interface: Spawning and Terminating the Web Server

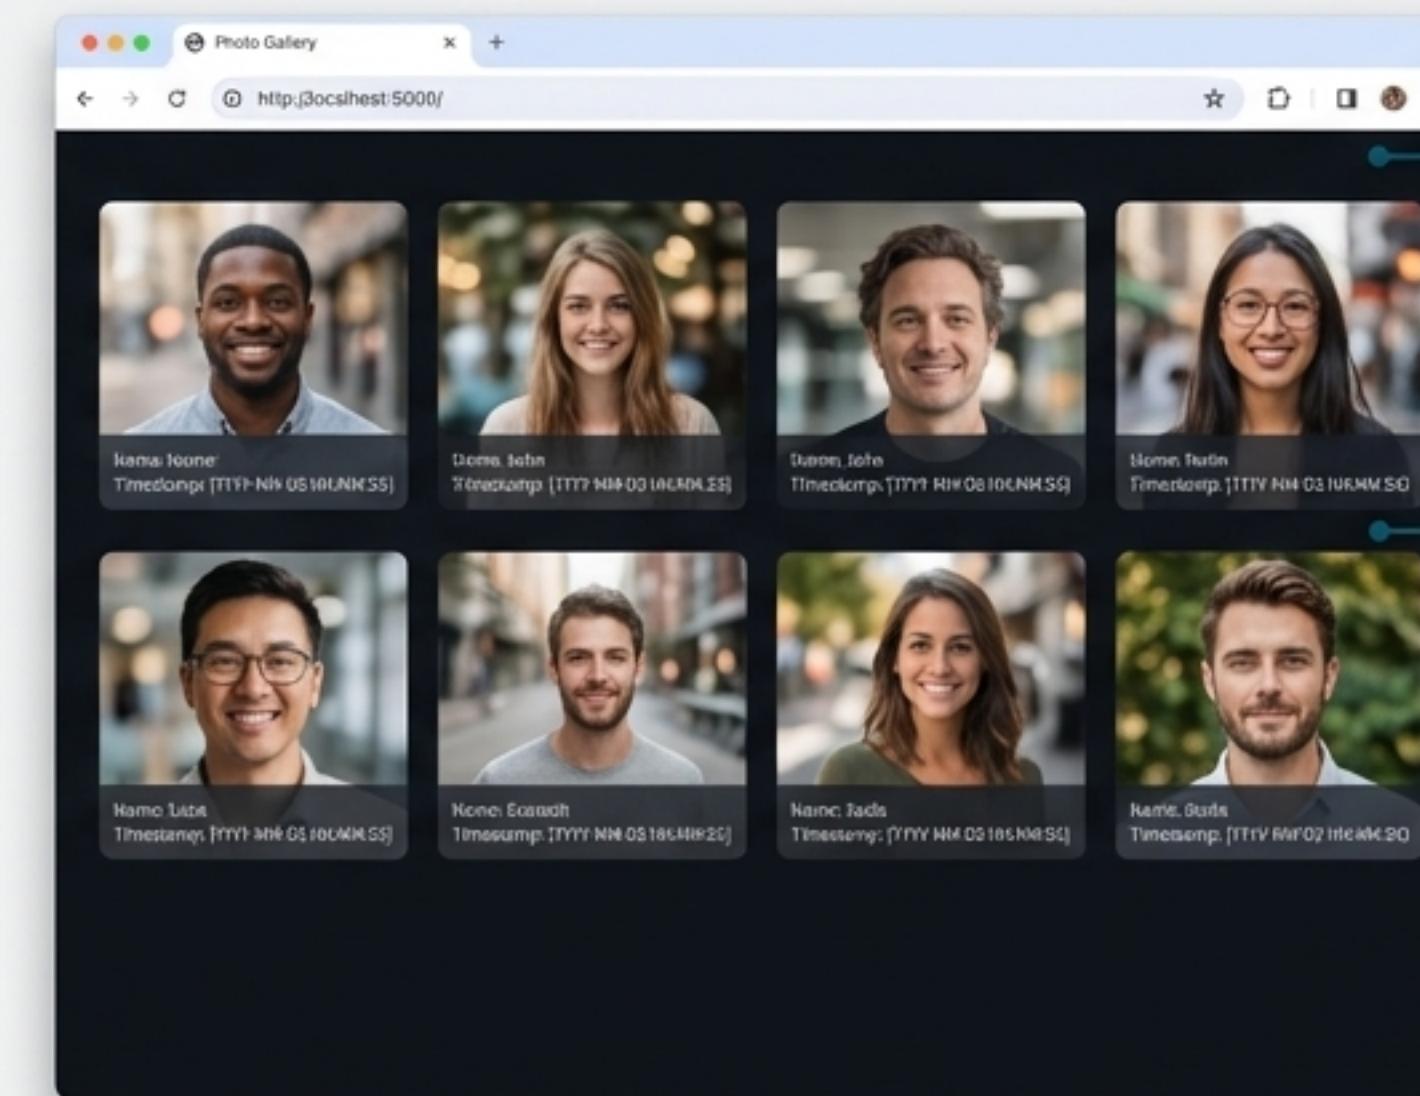


# The Public Face: A Simple and Efficient Web Interface

## Backend API (Flask)

- **GET /:**  
Serves the main HTML\_TEMPLATE.
- **GET /api/photos:**  
Scans the attendance\_photos/ directory, sorts images by modification time, and returns a clean JSON array of photo metadata (URL, name, timestamp).
- **GET /photos/<filename>:**  
Acts as a static file server to deliver the actual image assets.

## Frontend Logic (HTML/JS)



Uses modern CSS variables for a clean, dark-mode theme (--bg-color: #0d1117).

An async function, fetchPhotos(), polls the API. To prevent unnecessary DOM manipulation, a hash of current filenames is stored. The grid is only re-rendered if the hash changes, simulating a Virtual DOM concept for efficiency.

# Logic in Motion: Tracing an "Entry" Event



## 1. Recognition

A new face is detected and recognized as 'User X'.



## 2. State Check

The system checks if "User X" in present\_people: → **\*\*False\*\***



## 3. Action - Log

A new line is appended to 'lobby\_log.csv'.

`2023-10-27 10:00:00,ENTERED,User X,photos/Attendance\_UserX\_...jpg`



## 4. Action - Evidence

The current video frame is copied. The timestamp and name are burned onto the image using 'cv2.putText'. The result is saved to 'attendance\_photos/' via 'cv2.imwrite'.



## 5. State Update

The system updates its short-term memory.

present\_people["User X"] = <Current Unix Timestamp>

# Logic in Motion: Tracing an 'Exit' Event



## 1. Absence

For several consecutive frames, no faces are detected in the video stream.

## 2. State Iteration

The system loops through every name currently in the `present\_people` dictionary.

## 3. Timestamp Check

For 'User X', it calculates the elapsed time.

`Current Time - present\_people["User X"]`

Is the result > 3.0 seconds? —————→ **True**

## 4. Action - Log

A new line is appended to `lobby\_log.csv`.

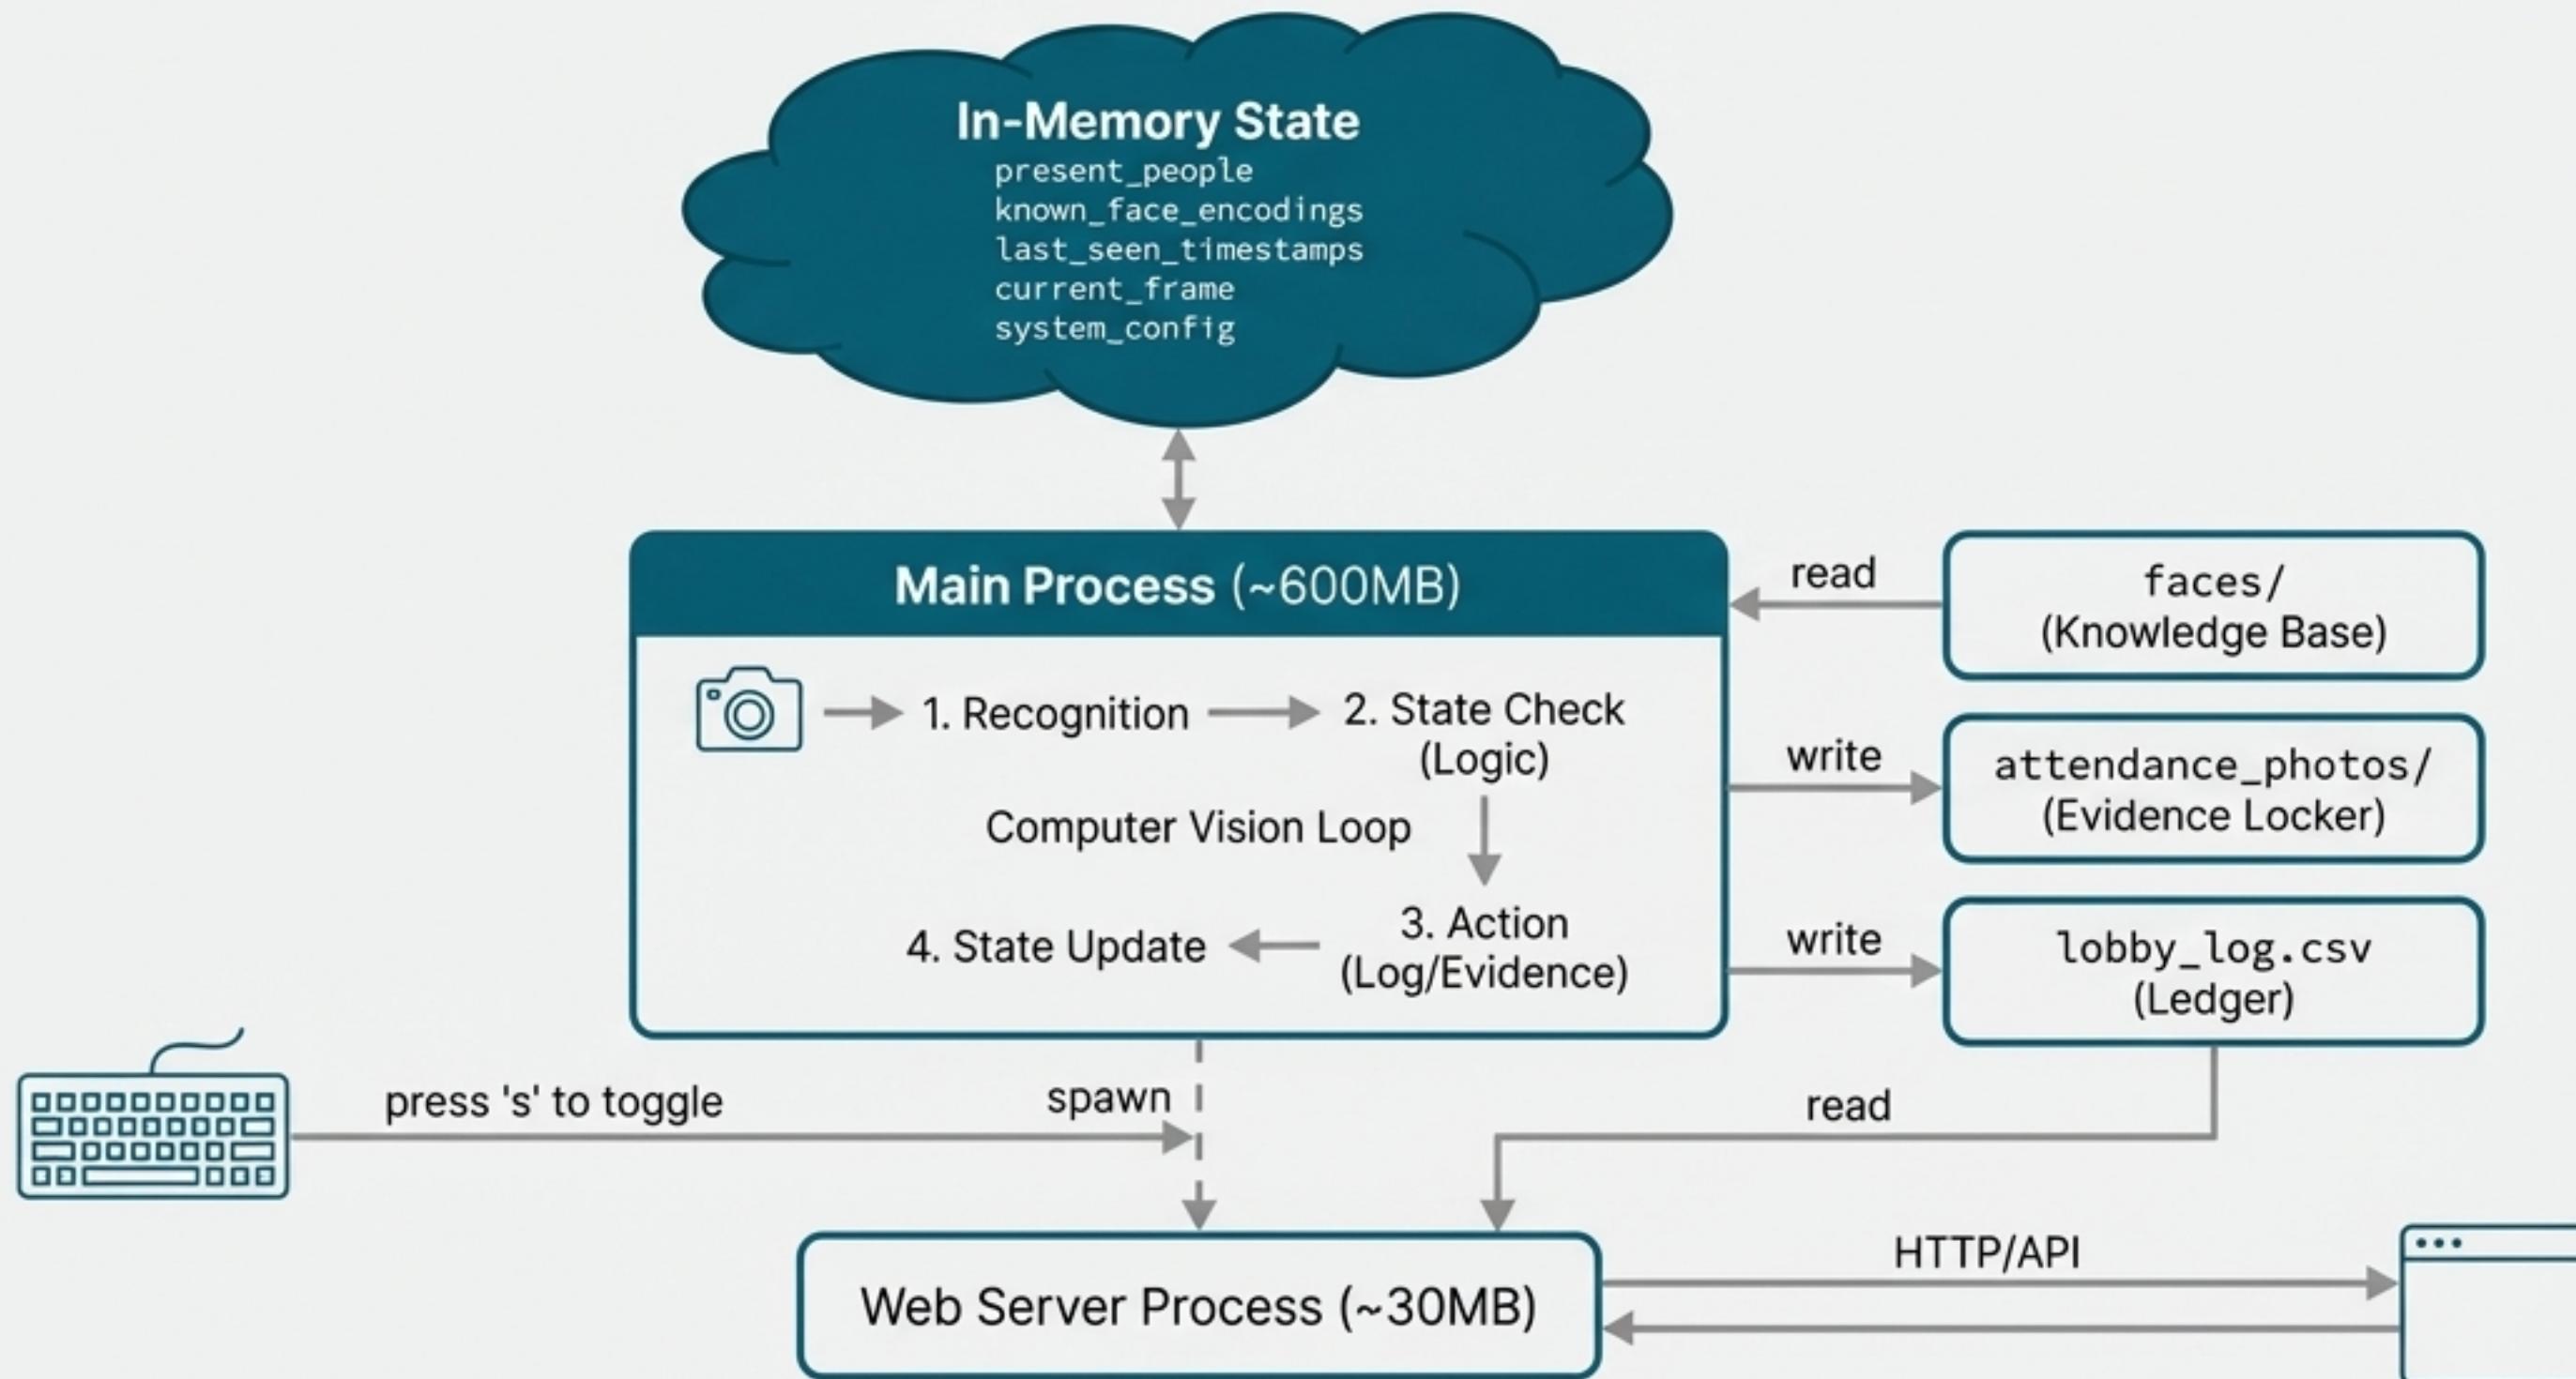
`2023-10-27 10:05:00, EXITED, User X`

## 5. State Update

The user is removed from short-term memory to allow for a new 'Entry' event later.

`del present\_people["User X"]`

# The Complete System Architecture



Attendance System v2.5 demonstrates a synthesis of efficient real-time processing, intelligent resource management, and stateful logic, creating a robust and performant monitoring tool from foundational components.